

Artificial Intelligence

Candidates for this certification have a foundational knowledge the procedures used to develop an artificial intelligence (AI) solution, as well as an understanding of the issues surrounding the governance, transparency, security, and ethics of AI. Successful candidates will be able to analyze and classify a problem. They should be able to demonstrate knowledge of data collection, data processing, and feature engineering strategies.

Candidates should be able to choose an appropriate algorithm for training a model, and understand the metrics used to evaluate model performance. They should understand the AI development lifecycle and how a production pipeline is used to allow for continuous improvement. Candidates at least 150 hours of instruction and/or exploration of artificial intelligence methodology and solutions.

1. AI Problem Definition

1.1 Identify the problem you are trying to solve using AI (e.g., user segmentation, improving customer service)

- Identify the need that will be addressed
- Find out what information comes in and what output is expected
- Determine whether AI is called for
- Consider upsides and downsides of AI in the situation
- Define measurable success
- Benchmark against domain or organization-specific risks to which the project may be susceptible

1.2 Classify the problem (e.g., regression, unsupervised learning)

- Examine available data (labeled or unlabeled?) and the problem
- Determine problem type (e.g., classification, regression, unsupervised, reinforcement)

1.3 Identify the areas of expertise needed to solve the problem

- Identify business expertise required
- Identify the need for domain (subject-matter) expertise on the problem
- Identify AI expertise needed
- Identify implementation expertise needed

1.4 Build a security plan

- Consider internal access levels or permissions
- Consider infrastructure security
- Assess the risk of using a certain model or potential attack surfaces (e.g., adversarial attacks on real-time learning model)

1.5 Ensure that AI is used appropriately

- Identify potential ways that the AI can mispredict or harm specific user groups
- Set guidelines for data gathering and use
- Set guidelines for algorithm selection from user perspective
- Consider how the subject of the data can interpret the results
- Consider out-of-context use of AI results



IT SPECIALIST EXAM OBJECTIVES

1.6 Choose transparency and validation activities

- Communicate intended purpose of data collection
- Decide who should see the results
- Review legal requirements specific to the industry with the problem being solved

2. Data Collection, Processing, and Engineering

2.1 Choose the way to collect data

- Determine type/characteristics of data needed
- Decide if there is an existing dataset or if you need to generate your own
- When generating your own dataset, decide whether collection can be automated or requires user input

2.2 Assess data quality

- Determine whether the dataset meets needs of task
- Look for missing or corrupt data elements

2.3 Ensure that data are representative

- Examine collection techniques for potential sources of bias
- Make sure the amount of data is enough to build an unbiased model

2.4 Identify resource requirements (e.g., computing, time complexity)

- Assess whether the problem is solvable with available computing resources
- Consider the budget of the project and the resources that are available

2.5 Convert data into suitable formats (e.g., numerical, image, time series)

- Convert data to binary (e.g., images become pixels)
- Convert computer data into features suitable for AI (e.g., sentences become tokens)

2.6 Select features for the AI model

- Determine which data features to include
- Build initial feature vectors for test/train dataset
- Consult with subject-matter experts to confirm feature selection

2.7 Engage in feature engineering

- Review features and determine what standard transformations are needed
- Create processed datasets

2.8 Identify training and test datasets

- Separate available data into training and test datasets
- Ensure test dataset is represented

2.9 Document data decisions

- List assumptions, predicates, and constraints upon which design choices have been reasoned
- Make this information available to regulators and end users who demand deep transparency



INFORMATION TECHNOLOGY
SPECIALIST

3. AI Algorithms and Models

3.1 Consider applicability of specific algorithms

- Evaluate AI algorithm families
- Decide which algorithms are suitable, e.g., neural network, classification (like decision tree, k means)

3.2 Train a model using the selected algorithm

- Train model for an algorithm with best-guess starting parameters.
- Tune the model by changing parameters
- Gather performance metrics for the model
- Iterate as needed

3.3 Select specific model after experimentation, avoiding overengineering

- Consider cost, speed, and other factors in evaluating models
- Determine whether selected model meets explainability requirements

3.4 Tell data stories

- Where feasible, create visualizations of the results
- Look for trends
- Verify that the visualization is useful for making a decision

3.5 Evaluate model performance (e.g., accuracy, precision)

- Check for overfitting, underfitting
- Generate metrics or KPIs
- Introduce new test data to cross-validate robustness, testing how model handles unforeseen data

3.6 Look for potential sources of bias in the algorithm

- Verify that inputs resemble training data
- Confirm that training data do not contain irrelevant correlations we do not want classifier to rely on
- Check for imbalances in data
- Guard against creating self-fulfilling prophecies based upon historical biases
- Check the explainability of the algorithm (e.g., feature importance in decision trees)

3.7 Evaluate model sensitivity

- Test for sensitivity of model
- Test for specificity of model

3.8 Confirm adherence to regulatory requirements, if any

- Evaluate outputs according to thresholds defined in requirements
- Document results

3.9 Obtain stakeholder approval

- Collect results and benchmark risks
- Hold sessions to evaluate solution



4. Application Integration and Deployment

4.1 Train customers on how to use product and what to expect from it

- Inform users of model limitations
- Inform users of intended model usage
- Share documentation
- Manage customer expectations

4.2 Plan to address potential challenges of models in production

- Understand the types of challenges you are likely to encounter
- Understand the indicators of challenges
- Understand how each type of challenge could be mitigated

4.3 Design a production pipeline, including application integration

- Create a pipeline (training, prediction) that can meet the product needs (may be different from the experiment)
- Find the solution that works with the existing data stores and connects to the application
- Build the connection between the AI and the application
- Build mechanism to gather user feedback
- Test accuracy of AI through application
- Test robustness of AI
- Test speed of AI
- Test application to fit size of use case (e.g., in AI for mobile applications)

4.4 Support the AI solution

- Document the functions within the AI solution to allow for maintenance (updates, fixing bugs, handling edge cases)
- Train a support team
- Implement a feedback mechanism
- Implement drift detector
- Implement ways to gather new data

5. Maintaining and Monitoring AI in Production

5.1 Engage in oversight

- Log application and model performance to facilitate security, debug, accountability, and audit
- Use robust monitoring systems
- Act upon alerts
- Observe the system over time in a variety of contexts to check for drift or degraded modes of operation
- Detect any way system fails to support new information

5.2 Assess business impact (key performance indicators)

- Track impact metrics to determine whether solution has solved the problem
- Compare previous metrics with new metrics when changes are made
- Act on unexpected metrics by finding problem and fixing it



IT SPECIALIST EXAM OBJECTIVES

5.3 Measure impacts on individuals and communities

- Analyze impact on specific subgroups
- Identify and mitigate issues
- Identify opportunities for optimization

5.4 Handle feedback from users

- Measure user satisfaction
- Assess whether users are confused (e.g., do they understand what the AI is supposed to do for them?)
- Incorporate feedback into future versions

5.5 Consider improvement or decommission on a regular basis

- Combine impact observations (e.g., business, community, technology trends) to assess AI value
- Decide whether to retrain AI, continue to use AI as is, or to decommission AI



INFORMATION TECHNOLOGY
SPECIALIST